

[Menu](#)[Browse Library](#) ▼[Training Library](#) > [CKAD Practice Exam: Configuration Solution Guide](#)

# CKAD Practice Exam: Configuration Solution Guide

Intermediate ⌚ 20m  [Bookmark](#)

## Check 1: Potential Solution

```
# Create a secret named app-secret in the yqe Namespace that stores key-
value pair of password=abnaoieb2073xsj
kubectl -n yqe create secret generic --from-
literal=password=abnaoieb2073xsj app-secret
# Create a Pod that consumes the app-secret Secret using a Volume that
mounts the Secret in the /etc/app directory. The Pod should be named app
and run a memcached container.
cat << EOF | kubectl apply -n yqe -f -
apiVersion: v1
kind: Pod
metadata:
  name: app
spec:
  containers:
  - name: app
    image: memcached
    volumeMounts:
    - name: secret
      mountPath: "/etc/app"
  volumes:
  - name: secret
    secret:
      secretName: app-secret
EOF
```

## Commentary

There are multiple types of Secrets but the most common is generic/opaque Secrets. This is the most likely type of Secret to appear on the exam. The Secrets documentation has a section linking to the different ways to create Secrets.

Using `kubectl` or manifest files are both reasonable choices but



using `kubectl`

[Skip to content](#)

Press  +  to open this menu [he](#)



The Secrets documentation also has [an example showing how to consume a Secret using a Volume](#). Customizing the example is the easiest way to achieve the result. Although the task doesn't specify which Namespace the Pod should be in, it must be in the same Namespace as the secret since secrets are a namespaced resource.

In the solution above, the file contents and `kubectl apply` command are combined into one. This approach allows you to use the exam notepad feature without any need to use command-line editors. You can then paste the command into the terminal to execute it. You can also create a file using `vim` and then paste in the example to modify. It is highly recommended to practice with `vim` before taking the exam.

## Content to review

- [Mastering Kubernetes Pod Configuration: Config Maps and Secrets - Storing and Accessing Sensitive Information Using Kubernetes Secrets](#)

## Suggested documentation bookmark(s)

- [Secrets](#)

## Check 2: Potential Solution

```
# Create a new ServiceAccount named secure-svc in the same namespace
where the deployment currently exists
kubectl create serviceaccount -n app secure-svc
# Create a file named patch.yaml and configure the replicas
serviceAccountName to be secure-svc
spec:
  template:
    spec:
      serviceAccountName: secure-svc
# Patch the existing Deployment
kubectl patch deployment -n app secapp --patch "$(cat patch.yaml)"
```

Comm

[Skip to content](#)

Press  +  to open this menu



Menu



Browse  
Library



ServiceAccount command. Deployments can be configured either before or after with custom service accounts for the replicas declared within. Deployments which have already been created can be patched to use a different ServiceAccount to the one that they were originally created with.

## Content to review

- [Configure Service Accounts for Pods](#)
- [Deployments](#)

## Suggested documentation bookmark(s)

- [Update API Objects in Place Using kubectl patch](#)

## Check 3: Potential Solution

```
# Start by generating a Pod manifest template specifying as many
parameters possible when using kubectl run - redirect and save the output
into a file named pod.yaml
(kubectl run --image=bash -n dnn secpod -l env=prod --dry-run=client -o
yaml -- /usr/local/bin/bash -c 'sleep 3600') > pod.yaml
# pod.yaml contains
apiVersion: v1
kind: Pod
metadata:
  creationTimestamp: null
  labels:
    env: prod
    name: secpod
    namespace: dnn
spec:
  containers:
  - args:
    - /usr/local/bin/bash
    - -c
    - sleep 3600
    image: bash
    name: secpod
    resources: {}
  dnsPolicy: ClusterFirst
  restartPolicy: Always
status: {}
# Use kubectl run to create the pod from the pod.yaml manifest.
```

[Skip to content](#)

Press **option** + **Q** to open this menu

the



Menu



Browse  
Library



```
labels:
  env: prod
  name: secpod
  namespace: dnn
spec:
  securityContext:
    fsGroup: 3000
  containers:
  - name: c1
    securityContext:
      runAsUser: 1000
    args:
    - /usr/local/bin/bash
    - -c
    - sleep 3600
    image: bash
    resources: {}
  - name: c2
    securityContext:
      runAsUser: 2000
    args:
    - /usr/local/bin/bash
    - -c
    - sleep 3600
    image: bash
    resources: {}
  dnsPolicy: ClusterFirst
  restartPolicy: Always
status: {}
# Using kubectl apply - create the pod resource within the cluster
kubectl apply -f pod.yaml
# Confirm that the pod has been successfully created
kubectl get pods -n dnn
# Confirm the the pod contains 2 containers c1 and c2 and that they have
been configured with the required security context
kubectl exec -it -n dnn secpod -c c1 -- id
uid=1000 gid=0(root) groups=3000
kubectl exec -it -n dnn secpod -c c2 -- id
uid=2000 gid=0(root) groups=3000
```

## Commentary

A Security Context defines privilege and access control settings either at the Pod layer or at the Container layer. To specify security settings for a Pod, include the `securityContext` field in the Pod spec section and/or in the Container section. The `runAsUser` field specifies the user id that all processes running within a container will run as.

Conte

[Skip to content](#)

Press **option** + **Q** to open this menu



Menu



Browse  
Library



- [kubect! Cheat Sheet \(--dry-run\)](#)

## Check 4: Potential Solution

```
# Start by generating a Pod manifest template specifying as many
parameters possible when using kubectl run - redirect and save the output
into a file named pod.yaml
(kubectl run web1 -n ca100 --image=nginx -l env=prod,type=processor --
port=80 --dry-run=client -o yaml) > pod.yaml
# pod.yaml contains
apiVersion: v1
kind: Pod
metadata:
  creationTimestamp: null
  labels:
    env: prod
    type: processor
  name: web1
  namespace: ca100
spec:
  containers:
  - image: nginx
    name: web1
    ports:
    - containerPort: 80
    resources: {}
  dnsPolicy: ClusterFirst
  restartPolicy: Always
status: {}
# Using vim open up and modify the pod.yaml template with the memory
request of 100Mi and a memory limit at 200Mi:
apiVersion: v1
kind: Pod
metadata:
  creationTimestamp: null
  labels:
    env: prod
    type: processor
  name: web1
  namespace: ca100
spec:
  containers:
  - image: nginx
    name: web1
    ports:
    - containerPort: 80
```

[Skip to content](#)

Press  +  to open this menu



Menu



Browse  
Library



```
# Using kubectl apply - create the pod resource within the cluster
kubectl apply -f pod.yaml
# Confirm that the pod has been successfully created
kubectl get pods -n ca100
# Confirm that the pod has the correct labels, port, limits, and requests
kubectl describe pod -n ca100 web1
```

## Content to review

- [Assign Memory Resources to Containers and Pods](#)
- [Pods](#)

## Suggested documentation bookmark(s)

- [kubectl Cheat Sheet \(run\)](#)

## Check 5: Potential Solution

```
# Create a new ConfigMap using the kubectl create command with the
required key/value pairs
kubectl create configmap config1 -n ca200 --from-literal COLOUR=red --
from-literal SPEED=fast
# Next, generate a Pod manifest template specifying as many parameters
possible when using kubectl run - redirect and save the output into a
file named pod.yaml
(kubectl run -n ca200 redfastcar --image=busybox --dry-run=client -o yaml
-- /bin/sh -c "env | grep -E 'COLOUR|SPEED'; sleep 3600") > pod.yaml
# pod.yaml contains:
apiVersion: v1
kind: Pod
metadata:
  creationTimestamp: null
  labels:
    run: car1
  name: car1
  namespace: ca200
spec:
  containers:
  - args:
    - /bin/sh
    - -c
    - env | grep -E 'COLOUR|SPEED'; sleep 3600
    image: busybox
    name: redfastcar
  dns:
  res:
  status: {}
```

Skip to content

Press **option** + **Q** to open this menu



Menu



Browse  
Library



```
kind: Pod
metadata:
  creationTimestamp: null
  labels:
    run: redfastcar
  name: redfastcar
  namespace: ca200
spec:
  containers:
  - args:
    - /bin/sh
    - -c
    - env | grep -E 'COLOUR|SPEED'; sleep 3600
    image: busybox
    name: redfastcar
    resources: {}
    envFrom:
    - configMapRef:
        name: config1
    dnsPolicy: ClusterFirst
    restartPolicy: Always
  status: {}
# Using kubectl apply - create the pod resource within the cluster
kubectl apply -f pod.yaml
# Confirm that the pod has launched successfully with the ConfigMap
managed env vars
kubectl logs -n ca200 redfastcar
COLOUR=red
SPEED=fast
```

## Commentary

ConfigMaps can be used to manage and maintain configuration key/value pairs which can then be later mounted into a Pod as environment variables.

## Content to review

- [ConfigMap as container environment variables](#)
- [Pods](#)

## Suggested documentation bookmark(s)

- [kubectl Cheat Sheet \(logs, run\)](#)

Skip to content

Press **option** + **Q** to open this menu





Menu



Browse  
Library



Mark as completed

Did you like this Resource?



Report an issue

## About the author



**Cloud Academy**

Instructor

Students

**2,617**

Labs

**83**

Courses

**9**

Learning paths

**18**

Digital skills are built at the intersection of knowledge, experience, and context. The fundamental building blocks of the training templates in our Library meet teams wherever they are along the cloud maturity curve, imparting the knowledge and experience needed to take them to the next level. Our training platform is intuitive and scalable. Most importantly, all training is easily customizable, which enables organizations to provide context and guidance for teams of any size. Teams leveraging Cloud Academy hit the ground running.

## Covered topics

Deployment

Compute

DevOps

Kubernetes



### ABOUT US

About

About

Skip to content

Press  +  to open this menu





Menu



Browse  
Library ▼



---

## COMMUNITY

[Join Discord Channel](#)

## HELP

[Help Center](#)

---

Copyright © 2024 Cloud Academy Inc. All rights reserved.

[Terms and Conditions](#)

[Privacy Policy](#)

[Sitemap](#)

[System Status](#)

[Manage your cookies](#)