



Menu



Browse
Library



Search in our library...

Deploy a Stateful Application in a Kubernetes Cluster

Deploying a Stateful Application in the Kubernetes Cluster

Lab Steps



Connecting to the Kubernetes Cluster



Inspecting the Kubernetes Cluster



Deploying a Stateful Application in the Kubernetes Cluster



Working with the Stateful Application



Monitoring Your Kubernetes Cluster Using Kubernetes Dashboard



Need help? Contact our support team

Here you can find the instructions for this specific Lab Step.

If you are ready for a real environment experience please start the Lab. Keep in mind that you'll need to start from the first step.

Start Lab

Introduction

Stateful applications are applications that have a memory of what happened in the past. Databases are an example of stateful applications. Kubernetes provides support for stateful applications through StatefulSets and related primitives.

This lab will deploy a replicated MySQL database as a StatefulSet. MySQL is one of the most popular databases in the world. This lab won't focus on the details specific to configuring MySQL. The focus is on the features of Kubernetes that allow a stateful application to be deployed. The Kubernetes community maintains a wide variety of stateful applications that are ready to deploy through [Helm](#). Helm acts as a package manager for Kubernetes and can deploy entire applications to a cluster using templates called charts. A list of charts is available [here](#). In addition to MySQL, there are charts for many other popular databases including MongoDB and PostgreSQL, as well as popular applications like WordPress and Joomla. The concepts illustrated in this lab will help you understand the common elements used to deploy all of these stateful applications in Kubernetes.

Managing your stateless and stateful applications with Kubernetes provides efficiencies and simplifies automation. However, before using StatefulSets for your own stateful applications, you should consider if any of th

apply:

Support



Skip to content

Press **option** + **Q** to open this menu e stateful



using specialized hardware and could effectively run on the same hardware used for stateless applications

- You value flexible reallocation of resources, consolidation, and automation over squeezing the most and having highly predictable performance

If any of the previous bullets apply to your situation, it may make sense to use Kubernetes for your stateful applications.

There are quite a few concepts at work in this lab step. Begin by reviewing the following background section and then jump into the instructions.

Background

ConfigMaps: A type of Kubernetes resource that is used to decouple configuration artifacts from image content to keep containerized applications portable. The configuration data is stored as key-value pairs.

Headless Service: A headless service is a Kubernetes service resource that won't load balance behind a single service IP. Instead, a headless service returns a list of DNS records that point directly to the pods that back the service. A headless service is defined by declaring the clusterIP property in a service spec and setting the value to None. StatefulSets currently require a headless service to identify pods in the cluster network.

Stateful Sets: Similar to Deployments in Kubernetes, StatefulSets manage the deployment and scaling of pods given a container spec. StatefulSets differ from Deployments in that the Pods in a stateful set are not interchangeable. Each pod in a StatefulSet has a persistent identifier that it maintains across any rescheduling. The pods in a StatefulSet are also ordered. This provides a guarantee that one pod can be created before following pods. In this lab, this is useful for ensuring the MySQL primary is provisioned first.

PersistentVolumes (PVs) and PersistentVolumeClaims (PVCs): PVs are Kubernetes resources that represent storage in the cluster. Unlike regular



MySQL replication: This lab uses a single primary, asynchronous replication scheme for MySQL. All database writes are handled by a single primary. The database replicas asynchronously synchronize with the primary. This means the primary will not wait for the data to be copied onto the replicas. This can improve the performance of the primary at the expense of having replicas that are not always exact copies of the primary. Many applications can tolerate slight differences in the data and are able to improve the performance of database read workloads by allowing clients to read from the replicas.

Instructions

1. In your SSH shell, enter the following to declare a ConfigMap to allow primary MySQL pods to be configured differently than replica pods:

[Copy code](#)

```
1 cat <<EOF > mysql-configmap.yaml
2 apiVersion: v1
3 kind: ConfigMap
4 metadata:
5   name: mysql
6   labels:
7     app: mysql
8 data:
9   master.cnf: |
10    # Apply this config only on the primary.
11    [mysqld]
12    log-bin
13   slave.cnf: |
14    # Apply this config only on replicas.
15    [mysqld]
16    super-read-only
17 EOF
```

This ConfigMap will be referenced later in the StatefulSet declaration. The `master.cnf` key maps to a value that declares a MySQL configuration which includes replication logs. The `slave.cnf` key maps to a MySQL configuration that enforces read-only behavior.

2. Create the ConfigMap resource:

[Copy code](#)

```
1 kubectl create -f mysql-configmap.yaml
```

[Copy code](#)

```
1 cat <<EOF > mysql-services.yaml
2 # Headless service for stable DNS entries of StatefulSet member
3 apiVersion: v1
4 kind: Service
5 metadata:
6   name: mysql
7   labels:
8     app: mysql
9 spec:
10   ports:
11     - name: mysql
12       port: 3306
13     clusterIP: None
14     selector:
15       app: mysql
16 ---
17 # Client service for connecting to any MySQL instance for reads
18 # For writes, you must instead connect to the primary: mysql-0.
19 apiVersion: v1
20 kind: Service
21 metadata:
22   name: mysql-read
23   labels:
24     app: mysql
25 spec:
26   ports:
27     - name: mysql
28       port: 3306
29     selector:
30       app: mysql
31 EOF
```

Two services are defined:

- A headless service (`clusterIP: None`) for pod DNS resolution. Because the service is named `mysql`, pods are accessible via `pod-name.mysql`.
- A service name `mysql-read` to connect to for database reads. This service uses the default `ServiceType` of `ClusterIP` which assigns an internal IP address that load balances request to all the pods labeled with `app: mysql`.

Database writes need to be sent to the primary. The primary is the first pod provisioned in the `StatefulSet` and assigned a name `mysql-0`. The pod is thus accessed by the DNS entry in the headless service for `mysql-0.mysql`.

4. Create the MySQL services:

[Copy code](#)

```
1 | kubectl create -f mysql-services.yaml
```



Menu



Browse
Library



 Copy code

```
1 cat <<EOF > mysql-storageclass.yaml
2 kind: StorageClass
3 apiVersion: storage.k8s.io/v1
4 metadata:
5   name: general
6   provisioner: kubernetes.io/aws-ebs
7   parameters:
8     type: gp2
9 EOF
```

The built-in aws-ebs storage provision is specified along with the type gp2.

6. Create the storage class:

 Copy code

```
1 | kubectl create -f mysql-storageclass.yaml
```

7. Enter the following command to declare the MySQL StatefulSet:

 Copy code

```
1 cat <<'EOF' > mysql-statefulset.yaml
2 apiVersion: apps/v1
3 kind: StatefulSet
4 metadata:
5   name: mysql
6 spec:
7   selector:
8     matchLabels:
9       app: mysql
10  serviceName: mysql
11  replicas: 3
12  template:
13    metadata:
14      labels:
15        app: mysql
16    spec:
17      initContainers:
18        - name: init-mysql
19          image: mysql:5.7.35
20          command:
21            - bash
22            - "-c"
23            - |
24              set -ex
25              # Generate mysql server-id from pod ordinal index.
26              [[ `hostname` =~ -([0-9]+)$ ]] || exit 1
27              ordinal=${BASH_REMATCH[1]}
28              echo [mysqld] > /mnt/conf.d/server-id.cnf
29              # Add an offset to avoid reserved server-id=0 value.
30              echo server-id=$((100 + $ordinal)) >> /mnt/conf.d/se
```

Skip to content

Press  +  to open this menu

conf.d/



Menu



Browse
Library



```
40     - name: config-map
41       mountPath: /mnt/config-map
42   - name: clone-mysql
43     image: gcr.io/google-samples/xtrabackup:1.0
44     command:
45     - bash
46     - "-c"
47     - |
48       set -ex
49       # Skip the clone if data already exists.
50       [[ -d /var/lib/mysql/mysql ]] && exit 0
51       # Skip the clone on primary (ordinal index 0).
52       [[ `hostname` =~ -([0-9]+)$ ]] || exit 1
53       ordinal=${BASH_REMATCH[1]}
54       [[ $ordinal -eq 0 ]] && exit 0
55       # Clone data from previous peer.
56       ncat --recv-only mysql-$((ordinal-1)).mysql 3307 |
57       # Prepare the backup.
58       xtrabackup --prepare --target-dir=/var/lib/mysql
59   volumeMounts:
60   - name: data
61     mountPath: /var/lib/mysql
62     subPath: mysql
63   - name: conf
64     mountPath: /etc/mysql/conf.d
65   containers:
66   - name: mysql
67     image: mysql:5.7
68     env:
69     - name: MYSQL_ALLOW_EMPTY_PASSWORD
70       value: "1"
71     ports:
72     - name: mysql
73       containerPort: 3306
74     volumeMounts:
75     - name: data
76       mountPath: /var/lib/mysql
77       subPath: mysql
78     - name: conf
79       mountPath: /etc/mysql/conf.d
80     resources:
81       requests:
82         cpu: 100m
83         memory: 200Mi
84     livenessProbe:
85       exec:
86         command: ["mysqladmin", "ping"]
87       initialDelaySeconds: 30
88       timeoutSeconds: 5
89     readinessProbe:
90       exec:
91         # Check we can execute queries over TCP (skip-networking)
92         command: ["mysql", "-h", "127.0.0.1", "-e", "SELECT 1"]
93       initialDelaySeconds: 5
94       timeoutSeconds: 1
95   - name: xtrabackup
96     image: gcr.io/google-samples/xtrabackup:1.0
97     ports:
98     - name: xtrabackup
99       containerPort: 3307
100     command:
101     - bash
102     - "-c"
103     - |
104       set -ex
105       cd /var/lib/mysql
```

Skip to content

Press **option** + **Q** to open this menu

data, if any.
n
tial "CHANGE M
ting replica.



```
115 [[ `cat xtrabackup_binlog_info` =~ ^(.*)[:space:
116 rm xtrabackup_binlog_info
117 echo "CHANGE MASTER TO MASTER_LOG_FILE='${BASH_REP
118 MASTER_LOG_POS=${BASH_REMATCH[2]}" > change_
119 fi
120 # Check if we need to complete a clone by starting i
121 if [[ -f change_master_to.sql.in ]]; then
122 echo "Waiting for mysql to be ready (accepting co
123 until mysql -h 127.0.0.1 -e "SELECT 1"; do sleep 1
124 echo "Initializing replication from clone positio
125 # In case of container restart, attempt this at-m
126 mv change_master_to.sql.in change_master_to.sql.or
127 mysql -h 127.0.0.1 <<EOF
128 $(<change_master_to.sql.orig),
129 MASTER_HOST='mysql-0.mysql',
130 MASTER_USER='root',
131 MASTER_PASSWORD='',
132 MASTER_CONNECT_RETRY=10;
133 START SLAVE;
134 EOF
135 fi
136 # Start a server to send backups when requested by p
137 exec ncat --listen --keep-open --send-only --max-cor
138 "xtrabackup --backup --slave-info --stream=xbstre
139 volumeMounts:
140 - name: data
141   mountPath: /var/lib/mysql
142   subPath: mysql
143 - name: conf
144   mountPath: /etc/mysql/conf.d
145 resources:
146   requests:
147     cpu: 100m
148     memory: 50Mi
149 volumes:
150 - name: conf
151   emptyDir: {}
152 - name: config-map
153   configMap:
154     name: mysql
155 volumeClaimTemplates:
156 - metadata:
157   name: data
158   spec:
159     accessModes: ["ReadWriteOnce"]
160     resources:
161       requests:
162         storage: 2Gi
163     storageClassName: general
164 EOF
```

There is a lot going on in the StatefulSet. Don't focus too much on the bash scripts that are performing MySQL-specific tasks. Some highlights to focus on, following the order they appear in the file are:

- `init-containers`: Run to completion before any containers in the Pod spec
 - `init-mysql`: Assigns a unique MySQL server ID starting from 100 for the first pod and incrementing by one, as well as copying the appropriate configuration file from the `config-map`. Note



files from the preceding pod. The `xtrabackup` tool performs the file cloning and persists the data on the data volume.

- `spec.containers`: Two containers in the pod
 - `mysql`: Runs the MySQL daemon and mounts the configuration in the `conf` volume and the data in the data volume
 - `xtrabackup`: A *sidecar* container that provides additional functionality to the `mysql` container. It starts a server to allow data cloning and begins replication on replicas using the cloned data files.
- `spec.volumes`: `conf` and `config-map` volumes are stored on the node's local disk. They are easily re-generated if a failure occurs and don't require PVs.
- `volumeClaimTemplates`: A template for each pod to create a PVC with. `ReadWriteOnce` accessMode allows the PV to be mounted by only one node at a time in read/write mode. The `storageClassName` references the AWS EBS gp2 storage class named `general` that you created earlier.

8. Create the StatefulSet and start watching the associated pods:

 [Copy code](#)

```
1 kubectl create -f mysql-statefulset.yaml
2 kubectl get pods -l app=mysql --watch
```

The `--watch` option causes any updates to the pods to be written to the output. It takes a few minutes to initialize all three replicas.

9. In order to view the logs in S3, open the AWS Management Console by clicking the following button to access the lab's cloud environment:



Open Cloud Environment

0%

Loading...

Average setup time: 2m 35s

10. Enter the following credentials created just for your lab session, and click **Sign In**:

The 2GiB PVs are listed here as each pod is created. Notice the **Tags** which relay information about the PV and associated PVC.

	Name	Volume ID	Type	Size	IOPS	Throughput	Snapshot
<input checked="" type="checkbox"/>	-	vol-0b3f979803c717e19	gp2	2 GiB	100	-	-
<input type="checkbox"/>	-	vol-0799a5b9aa9b51583	gp2	2 GiB	100	-	-
<input type="checkbox"/>	-	vol-0799a5b9aa9b51583	gp2	2 GiB	100	-	-

Tags

Key	Value
kubernetes.io/created-for/pvc/namespace	default
ca-environment-session-uuid	705294fe-1ad8-41ce-aa7e-dfc1525494ed
ebs.csi.aws.com/cluster	true
ca-scope	lab
ca-persistent	false
ca-environment	production
CSIVolumeName	pvc-1f0eaff1-191d-49ab-9c03-cbad409bf5b1
kubernetes.io/created-for/pv/name	pvc-1f0eaff1-191d-49ab-9c03-cbad409bf5b1
ca-environment-session-id	1424187
kubernetes.io/created-for/pvc/name	data-mysql-1

12. Return to your SSH shell and press ctrl+C to stop the watch when you see both containers (2/2) in the **mysql-2** pod running:

```

NAME          READY    STATUS             RESTARTS   AGE
mysql-0       0/2      Init:0/2           0           3s
mysql-0       0/2      Init:1/2           0          43s
mysql-0       0/2      PodInitializing    0          56s
mysql-0       1/2      Running            0          57s
mysql-0       2/2      Running            0          1m
mysql-1       0/2      Pending            0           0s
mysql-1       0/2      Pending            0           0s
mysql-1       0/2      Pending            0           2s
mysql-1       0/2      Init:0/2           0           2s
mysql-1       0/2      Init:1/2           0          42s
mysql-1       0/2      Init:1/2           0          55s
mysql-1       0/2      PodInitializing    0          1m
mysql-1       1/2      Running            0           1m
mysql-1       2/2      Running            0           1m
mysql-2       0/2      Pending            0           0s
mysql-2       0/2      Pending            0           0s
mysql-2       0/2      Pending            0           1s
mysql-2       0/2      Init:0/2           0           1s
mysql-2       0/2      Init:1/2           0          43s
mysql-2       0/2      Init:1/2           0          54s
mysql-2       0/2      PodInitializing    0          1m
mysql-2       1/2      Pending            0           1m
mysql-2       2/2      Running            0           1m

```



Menu



Browse
Library



 Copy code

```
1 | kubectl describe pv
2 | kubectl describe pvc
```

The PV descriptions include the AWS **VolumeIDs**, file system types (**FSType**), and associated PVC (**Claim**). The PVC description includes whether the PVC is currently **Bound** to a pod.

14. Get the StatefulSet to confirm the current number of replicas matches the desired:

 Copy code

```
1 | kubectl get statefulset
```

Summary

In this lab step, you created several Kubernetes cluster resources to deploy the MySQL database as an example stateful application:

- A ConfigMap for decoupling primary and replica configuration from the containers
- Two Services: one headless service to manage network identity of pods in the StatefulSet, and one to load balance read access to the MySQL replicas
- A StorageClass to provision EBS PVs dynamically
- A StatefulSet that declared two init-containers, two containers, and one PVC template

You observed the ordered sequence of pods being initialized and the PVs created in AWS to facilitate the StatefulSet.



ABOUT US

Skip to content

Press  +  to open this menu



Menu



Browse
Library ▼



About Circus Street

COMMUNITY

Join Discord Channel

HELP

Help Center

Copyright © 2024 Cloud Academy Inc. All rights reserved.

[Terms and Conditions](#)

[Privacy Policy](#)

[Sitemap](#)

[System Status](#)

[Manage your cookies](#)